

DUG — Revenue Model & Embedded Products Strategy

Evergreen, Flexible Revenue Architecture · May 2026 · Internal

Executive Summary

DUG's revenue model must work at every order of magnitude simultaneously: a homeowner paying \$25 for a renewal sanity check, a risk manager paying \$400 for a pre-broker consult, an MGA paying \$15,000 for a program audit, and a Tier-1 carrier paying \$500,000/year for a dedicated underwriting review bench with SLA guarantees. The platform fee structure, engagement tiers, and contract mechanics are different at each level — but they share a single underlying marketplace engine.

Alongside transaction fees, DUG has a high-margin embedded insurance distribution opportunity: as underwriters scale their consulting income through the platform, they need E&O coverage. DUG can own that distribution channel — embedding E&O purchasing directly into the underwriter onboarding and income milestone flow, earning commission on every policy sold.

Part 1 — Engagement Tiers & Fee Architecture

Four tiers, one platform engine. Fee structures vary by tier but all flow through the same transaction and payment infrastructure.

Tier	Who & What	Typical Range	DUG Fee Model
Consumer (Tier 1)	Homeowners, personal auto insureds, individuals challenging a renewal or denial	\$25 – \$500 flat \$50–150/hr	20% platform fee Min. \$5 per transaction
Professional (Tier 2)	Risk managers, commercial insureds, SMB owners, brokers ordering second opinions	\$250 – \$5,000 flat \$100–250/hr	15% platform fee Capped at \$750/transaction
Enterprise (Tier 3)	MGAs, mid-market carriers, AI companies, reinsurers; SLA-backed response times	\$5,000 – \$100,000 Contract-based	10% platform fee + \$500 contract setup SLA tier: +\$1,000–5,000/mo
Strategic (Tier 4)	Tier-1 carriers, national programs; dedicated underwriting bench; annual contracts	\$100,000 – \$1M+ Annual retainer	Negotiated: 5–8% + annual platform license \$25,000–100,000/yr base

Fee Mechanics by Tier

Tier 1 — Consumer

- Fixed percentage (20%) applied at payment. No negotiation, no exceptions. Minimum \$5 floor ensures platform sustainability on micro-transactions.
- Payment processed at engagement acceptance. DUG holds funds and releases to underwriter upon delivery confirmation (or 72-hour auto-release).
- No SLA. Response expectation set by requester at posting.

Tier 2 — Professional

- 15% fee, capped at \$750 per transaction. This cap protects high-value professional relationships and encourages larger engagements.
- Milestone payments available: 50% on acceptance, 50% on delivery. DUG arbitrates disputes.
- Requester can post SLA preference; underwriter accepts or declines. No platform-level SLA guarantee.

Tier 3 — Enterprise

- 10% fee on transaction value plus a \$500 contract setup fee per engagement series.
- SLA options: 24-hour (\$1,000/mo add-on), 48-hour (included), 72-hour (default). SLA breach refunds setup fee.
- Dedicated account manager assigned at \$20,000+ annual volume. Quarterly business reviews.
- Enterprise clients can build a curated bench of preferred underwriters — retained pool, first right of refusal on availability.

Tier 4 — Strategic

- Annual contract, negotiated. 5–8% on transaction volume + annual platform license (\$25K–\$100K depending on volume and services).
- White-glove onboarding for carrier underwriting teams. Custom SLA definitions, escalation paths, and quality rubrics.
- DUG provides aggregate analytics: turnaround time, accuracy scoring, specialty depth. Delivered quarterly.
-

Extended Budget Models: Bounty & Milestone

Beyond flat, hourly, and volunteer, two additional budget models unlock the lowest and highest margin work DUG can support. Both are built on the same transaction infrastructure — no new payment rails required.

Per-Find / Bounty (Audit Swarm Model)

A carrier or MGA uploads a batch of submissions — 500 renewals, 2,000 submissions, a full book segment — and defines what a “find” looks like: a mispriced risk above a threshold, a wrong class code, a missing endorsement, an undisclosed exposure, a rating factor inconsistency. The community works through the batch at \$5–10 per confirmed find. A second-reviewer confirmation layer prevents false-positive gaming. DUG takes its platform percentage off each confirmed payout.

- **Job type:** portfolio_audit. Available to carrier, MGA, and reinsurer requesters.
- **Gamification:** Accuracy multipliers (high confirmation rate = higher per-find earnings), streak bonuses, category leaderboards, and “gold find” bonuses for defects above a dollar threshold. Every audit completion contributes to Dojo record.

- **Economics for carriers:** A book with 10,000 submissions where 3% carry material defects representing \$50K+ in combined premium correction pays for itself at \$8/find many times over. Cheaper, faster, and more thorough than internal audit teams.
- **Platform data:** Stored as find_bounty_cents on the job record. Total bounty pool (budget_cents) is an optional cap. DUG's percentage applies to each confirmed find payout individually.

Milestone Billing (High-Value Engagements)

Program design, portfolio restructuring, and complex multi-week engagements do not fit flat or hourly billing. A \$1,000/hr senior underwriter working a 6-week program design engagement needs milestone-based payment: funds released at defined deliverable gates, not on an hourly clock. DUG holds funds in escrow and releases per milestone confirmation. The platform's job at this tier is to hold money and stay out of the room.

- **Structure:** Up to 5 named milestones per engagement (e.g., "Kickoff & scoping" 20%, "Draft framework" 30%, "Final program design" 50%). Each milestone has a label and a dollar amount. Total is the engagement value against which DUG's fee applies.
- **Rate:** No rate cap. \$1,000/hr is an appropriate market rate for senior program design work. The platform does not price-govern the high end — the market sets it and DUG takes its percentage.
- **E&O trigger:** Program design engagements above \$10,000 total value surface a strong recommendation (not a block) to ensure Standard or Professional E&O tier is active. The larger the engagement, the larger the tail risk for the underwriter.

Platform Philosophy: Leakage Is a Feature, Not a Bug

DUG will not invest in controls designed to prevent requesters from taking underwriter relationships off-platform. A carrier that hires a DUG underwriter for a full-time role — or routes a repeat engagement directly — is a DUG success story, not a defection. That underwriter built their reputation through the Dojo, was found through the marketplace, and demonstrated enough value to be worth stealing. The next enterprise client will find someone else on DUG. The community grows either way.

The platform's value at high-value tiers is what naturally retains engagements: E&O coverage the underwriter cannot easily replicate independently, milestone escrow that protects both parties, documented scope of work that limits reliance claims, and the reputation system that made the right person findable in the first place. If the work is good enough that people are poaching DUG users, that is the right problem to have. Stay focused on expanding the community and delivering further value to the industry.

IP ownership of all analyses remains with the engaging party. DUG retains anonymized, aggregated performance data.

Part 2 — Embedded Insurance Products

The embedded E&O play is the highest-margin revenue line DUG can build over a 3-year horizon. It requires no additional infrastructure beyond what already exists — the underwriter profile, the income tracking, and the trust/compliance framework. What it requires is an MGA or carrier partner willing to write a group master E&O policy distributed through DUG.

The Opportunity

Every underwriter doing paid work through DUG's marketplace needs professional liability (E&O) coverage. Right now they either: (a) carry it from a prior employer and are burning that goodwill, (b) bought it expensively as a sole proprietor, or (c) don't have it and are attesting falsely. None of these outcomes are good for DUG or for the underwriter.

DUG's position: we are the distribution channel for a homogeneous, understandable risk class — independent underwriters doing advisory work. We know exactly what they earn, what lines they work, how many transactions they complete, and what their peer reputation score is. That data makes pricing trivial for a carrier partner. The group master structure means individual underwriters get coverage at group rates (40–60% cheaper than individual policies) and DUG earns 10–15% commission.

Product Design

Group Master E&O Policy

DUG is the named insured on a master policy. Marketplace underwriters are added as covered parties via a certificate of participation. Coverage attaches per-engagement or annually based on tier. Target carriers: Hiscox, Markel, ProAssurance — all have appetite for this class.

Trigger: Income Milestones

The E&O prompt appears automatically at three points in the underwriter journey: (1) marketplace activation, (2) first paid engagement completed, (3) \$1,000 in cumulative earnings. At each trigger, the underwriter sees their current coverage status and a one-click enrollment option.

Pricing Structure

Three tiers: Basic (\$199/yr, covers up to \$50K in annual DUG earnings), Standard (\$399/yr, up to \$150K), Professional (\$699/yr, unlimited DUG earnings). All include \$1M per occurrence / \$2M aggregate. DUG earns 12% commission on premium — approximately \$24–84 per policy per year.

Cyber Add-On

At Standard and Professional tiers, a cyber endorsement is available (\$99–149/yr add-on). Underwriters handling sensitive submission data (carrier, commercial insured) have meaningful cyber exposure. This is an easy upsell with the same carrier partner.

Revenue Projection

At 500 active marketplace underwriters with 60% E&O adoption: ~\$60K–120K/yr in commission revenue. At 2,000 underwriters: \$240K–480K/yr. This is pure margin — no infrastructure cost beyond the enrollment flow.

Part 3 — Underwriting Assessment & Interview Platform

A standalone B2B revenue line built on the same Dojo infrastructure — sold directly to carriers as a hiring and screening product.

The Market Gap

The developer world has CoderPad (live coding interviews with real-time observation), HackerRank (async scored challenges), and Karat (interviews-as-a-service). Insurance has none of this. Carriers still screen underwriting candidates with HR phone calls, subjective panel interviews, and the occasional PDF case study emailed the night before — no standardization, no scoring, no benchmarking against any external reference point.

General pre-employment platforms (Criteria Corp, Vervoe, TestGorilla) have loose insurance templates but nothing built around actual underwriting judgment. No one owns this space. DUG's Dojo infrastructure already does 80% of what the product requires.

Product 1 — Live Interview Sandbox

A timed, structured environment where a candidate works a real underwriting case while the hiring manager observes in a live side-panel — same concept as CoderPad for coding interviews. The case does the screening; the interviewer asks follow-up questions mid-session. After the session, the platform auto-generates a scorecard comparing the candidate's conclusions against the DUG community benchmark for that specific case.

- **Candidate experience:** Works a live risk case inside the Dojo interface. Timer runs. Notes, risk conclusions, and premium rationale are captured in real time.
- **Interviewer experience:** Separate observer panel shows candidate work in real time. Can inject mid-session prompts ("What if the occupancy changes?" / "The broker just disclosed a prior claim.").
- **Post-session output:** Auto-generated scorecard comparing candidate conclusions against the DUG community benchmark for that case. Carrier gets a percentile ranking, not just a gut feeling.
- **Async mode:** Candidate completes case independently; carrier reviews scored results on their own timeline. Same infrastructure, no scheduling coordination required.

Product 2 — Assessment Library

A catalog of 200+ cases across lines — GL, property, E&O, cyber, construction, marine, ag, D&O — that carriers license by the seat or by the year. Different difficulty levels for junior vs. senior screening. DUG maintains and updates the library. Carriers use it out of the box or build bespoke cases on top of the platform infrastructure.

- **Standard library license:** Access to full case library. Per-assessment pricing (\$50–150/use) or annual flat (\$5,000–15,000/yr by volume).
- **Custom case development:** DUG builds bespoke cases reflecting a carrier's specific book, appetite, or risk culture. One-time development fee (\$2,000–10,000) plus library hosting.
- **White-label:** Large carriers can license the platform under their own brand for internal training and hiring. Annual contract, \$25,000–75,000/yr depending on seat count.

- **Candidate-facing (prep product):** DUG Dojo Premium — candidates pay to practice on carrier-specific case packs before their interview. Analogous to LeetCode Premium. \$19–49/mo.

The Benchmark Moat

HackerRank can score a coding problem because there is a correct answer. Underwriting judgment is more nuanced — you need a large community of credentialed practitioners who have already worked those same cases to generate a meaningful distribution. DUG’s Dojo community is that benchmark. No competitor can replicate it without years of community building. The more the community grows, the more defensible and precise the scoring becomes. This is the core moat of the assessment product.

Critically, this product is independent of transaction fee revenue. It scales as a software line — growing with carrier seat counts and hiring volume rather than engagement GMV. That independence makes it a valuable hedge in the overall revenue portfolio.

Revenue Model

- **Per-assessment:** \$50–150 per live or async session. No commitment required — lowest friction entry for carriers trialing the product.
- **Annual library license:** \$5,000–15,000/yr for unlimited assessments up to seat count. Recurring, predictable SaaS revenue independent of marketplace GMV.
- **Custom development + white-label:** \$25,000–75,000/yr. Strategic accounts with dedicated infrastructure and carrier branding.
- **Dojo Premium (candidate-side):** \$19–49/mo subscription. Candidates prep for carrier interviews using line-specific case packs. Flywheel effect: more Dojo Premium users = richer benchmark data = stronger carrier product.

Build Roadmap

- **Phase 1 (now — Month 6):** Dojo practice cases are the foundation. Build the benchmark dataset. Every Dojo completion enriches the scoring distribution used by the carrier product later.
- **Phase 2 (Month 6–12):** Async assessment mode live. First carrier beta pilots at per-assessment pricing. Identify 3–5 carrier design partners willing to trial the product in exchange for input on the scoring rubric.
- **Phase 3 (Month 12–24):** Live interview sandbox mode with real-time observer panel. Annual library licenses. Dojo Premium launched as candidate-side prep product. ATS integrations (Workday, Greenhouse, Lever) to embed assessments directly into carrier hiring workflows.
- **Phase 4 (Month 24+):** White-label offering for Tier-1 carriers. Certification org partnerships (The Institutes, National Alliance) where Dojo Premium cases count toward CE credit — cross-pollinating the assessment and partnership revenue lines.

Part 4 — Revenue Source Tracking Architecture

As the community grows, new revenue sources will emerge that we can’t fully anticipate today. The platform tracks every transaction against a source tag from day one, so we can identify breakout categories before they become obvious.

Source Tags (Built Into Every Transaction)

requester_type — Primary segmentation — the who

carrier | mga | reinsurer | broker | agent | risk_manager | insured_commercial | insured_personal | tech_ai | other

engagement_tier — The size bucket

consumer | professional | enterprise | strategic

job_type — The what — tracks emerging use cases

renewal_review | pre_broker_consult | ai_benchmark | coverage_dispute | program_design | ...

acquisition_source — How the requester found us

organic | referral | enterprise_contract | dojo_graduate | community_member | paid_acquisition

underwriter_source — Which onboarding track the fulfilling underwriter used

credential_path | dojo_path

embedded_product — Whether an embedded product was also purchased in this cycle

eo_basic | eo_standard | eo_professional | cyber_addon | none

Metrics to Watch as Community Scales

- **GMV by tier:** Are consumer transactions actually converting, or is volume concentrated in Tier 2/3? Informs pricing floor decisions.
 - **Job type breakout:** Watch for emerging job types that don't fit the current taxonomy — those are new product lines.
 - **E&O attach rate by income tier:** If high-earning underwriters aren't buying E&O through DUG, find out why. It's either price or process friction.
 - **Enterprise contract conversion:** What % of Tier 3 requesters convert to annual contracts? This is the LTV inflection point.
 - **Dojo-to-marketplace conversion:** How many Dojo members activate as paid underwriters? This is the community flywheel metric.
 - **Requester repeat rate:** What % of requesters come back for a second job? High repeat = product-market fit. Low repeat = quality problem.
 - **Underwriter utilization rate:** Jobs completed / underwriters active. Low utilization means supply > demand — focus on demand side.
-

Part 5 — Revenue Roadmap by Stage

Now — Community Launch (\$0 GMV target)

- No revenue yet. Track everything from day one anyway — source tags cost nothing to implement and are invaluable retrospectively.
- Identify the first 10 underwriters who would do paid work if the marketplace existed. These are your beta Tier 2 clients.
- Begin MGA conversations for the group E&O master policy. This takes 3–6 months to bind — start now.

Marketplace Beta (Month 3–6) (\$0 – \$50K GMV GMV target)

- Consumer and Tier 2 transactions live. 20%/15% fee model active.
- E&O enrollment prompt live at marketplace activation. Even if the master policy isn't bound yet, capture intent and email.
- First enterprise conversations based on who is posting Tier 2 jobs. Upgrade path: 'You've posted 5 requests — want SLA guarantees and a preferred bench?'

Growth (Month 6–18) (\$50K – \$1M GMV GMV target)

- Enterprise tier (Tier 3) live with SLA mechanics and account managers.
- Group E&O master policy bound. E&O enrollment becomes a revenue line.
- First annual enterprise contracts signed. Track contract value separately from transactional GMV.
- Cyber add-on available. Dojo performance scores visible on underwriter profiles — credentialing becomes a competitive feature.

Scale (Month 18–36) (\$1M – \$10M GMV GMV target)

- Strategic tier (Tier 4) active. First Tier-1 carrier under contract.
 - E&O renewal cohort producing recurring commission revenue without acquisition cost.
 - Job type analytics surfacing new use cases. Consider launching purpose-built products (e.g., dedicated AI benchmarking subscription).
 - Consider captive or MGA structure for E&O distribution — at scale, the economics shift dramatically in favor of owning the paper.
-